

elife-claim-trees

Toward claim-level scientific knowledge

Zach Mainen — Mainen Lab, Champalimaud Foundation

2026-05-11

Docs: <https://zmainen.github.io/elife-claim-trees/docs/>

How does science change with AI?

One answer: **the unit of knowledge changes.**

Science currently works with **papers** as the fundamental unit. But papers are coarse — each one contains dozens of individual elements: a statistical result, a figure panel showing a correlation, a hypothesis, a methodological decision.

We call these elements **claims**.

AI creates the opportunity to work at the level of individual claims — to extract them, structure them, link them across papers, and eventually to compose science in claim-level units from the start.

What is a claim?

A claim is **one declarative proposition** from a paper, tied to its evidence.

```
claim: Doubling distal dendritic inhibition reduces  
       somatic firing rate from ~5.5 Hz to ~0.2 Hz.  
panel: fig4, fig5  
role:  empirical
```

- Anchored to a **specific figure panel** — where the evidence lives
- Classified by **role** — empirical result, hypothesis, prediction, control
- Linked to **other claims** — supports, requires, dissociates-with

Claims need not be stated explicitly. Hypotheses and predictions are often **implicit** in the structure of the experiment, not written out in the text.

"The unit of publication is not a figure or a statistic — it is a computer program." — Kenneth Harris, formulating this idea at an early eLife meeting

A claim is that computation: data in, analysis, result out. The panel is where one computation lives.

The vision: a claim-level knowledge graph

Today we have the **paper-citation graph**: "Paper A cites Paper B." Useful, but coarse.

A **claim-level knowledge graph** would be much finer-grained — describing the relationships of dependency between individual claims across a corpus of papers in a field.

- At **N = 10 papers**, it's a demonstration
- At **N = 1,000**, it's a research instrument
- At **N = 100,000**, it's the structure of a field

Unlike the paper-citation graph, you can ask: *"what does the field actually say about X?"* and get an answer grounded in specific, traceable propositions.

Two paths to get there

Path 1: Extract from existing papers. Scientists don't yet work in claim-level units. For the existing corpus — and for papers that will be submitted before anyone adopts claim-native authoring — we can do a conversion: interpret a paper into a structured set of claims using AI agents. **This is what we've built.**

Path 2: Author tools. Help scientists compose submissions as claim-units from the start — making the paper composition process more claim-like. Reduces the need for extraction over time. **Not yet built, but part of the plan.**

Both paths serve the same goal. The extraction pipeline (Path 1) is what's on the table today.

Why this matters beyond extraction

Claim-level peer review. QED Science has been pioneering this. When a paper is structured as claims, reviewers can comment on specific propositions, disagreements localize, and the process becomes more transparent.

Benefits for authors. Thinking in claims helps authors clarify what they're actually arguing — which results are empirical, which are interpretive, which depend on which.

Agent-driven science. A longer-horizon experiment: a sandbox where AI agents publish at claim-level granularity. Claims are the right unit for this — papers are too coarse for agents to compose and review at scale.

What we've built: a working draft

I've developed an extraction pipeline using language models as agents who review the paper — from the figures and the text — and extract a series of claims.

- **Working** — tested on 10 eLife papers, with measured results
- **A draft** — it needs investigation and improvements
- **Incorporable** — could be added to a website as an optional feature

Two deployment modes:

- **Author-approved:** an optional add-on the author reviews before publication
- **Public-domain:** run on published open-access papers (eLife, bioRxiv) without author involvement

The rest of this talk: how the pipeline works, how well it works, and where eLife could plug in.

Why not just ask an LLM to read the paper?

We tried. A single model reading the whole paper fails in predictable ways:

- **Quantitative hallucination** — invents precision ("~5 Hz" becomes "5.2 +/- 0.3 Hz")
- **Wrong panel assignment** — attaches claims to the wrong figure panels
- **Misses methodology** — scope conditions, controls, analytical decisions are flattened
- **Misses deductive structure** — implicit hypotheses and predictions disappear

Each failure mode requires its own architectural fix.

The pipeline



What each agent catches

AGENT	READS	GOOD AT
Results-reader	Abstract + results text	Hypotheses, predictions, synthesis claims
Caption-reader	Figure captions, panel by panel	Exact numbers, panel anchoring
Structure-reader	Methods + supplements	Controls, scope conditions, methodological claims

No agent sees another's output. The reconciler (a separate model call) aligns their outputs semantically and flags disagreements.

The external reviewer

Many important claims are **never stated explicitly** in the paper. The hypothesis the paper tests, the predictions that follow from its model — these live in the architecture of the experiment, not in the prose.

Human curators infer these. A prose-reading LLM misses them.

The external reviewer is a separate model pass that reads the whole paper plus the reconciled draft, looking for:

- Implicit hypotheses
- Predictions deduced from the modelling structure
- Claims that function as controls (not just results)

This single addition lifts role classification from **~65% to ~83%** agreement with curators.

The full pipeline

STEP	WHAT HAPPENS	WHO DOES IT
1. Prepare	Fetch PDF, slice into sections	Automated
2 - 3. Extract	Three independent agents read their slices	Sonnet (x3)
4. Reconcile	Merge three lists into one draft	Opus
4.5. Review	Recover implicit structure	Opus
5. Gate	Human reviews or Opus substitutes	Configurable
6 - 7. Write	Produce claim files with frontmatter	Automated

Input: one DOI. **Output:** 50-90 structured claim files.

Per paper: ~\$7 API cost, ~15 minutes.

Claim roles — the 9 rhetorical functions

Every claim has a **role** — what function it serves in the paper's argument. The methodology defines 9:

ROLE	WHAT IT MARKS
hypothesis	The paper's central bet — often implicit
prediction	A testable consequence of the hypothesis
empirical	A measured or computed result, panel-grounded
control	An empirical result that rules out an alternative
scope	A boundary condition on the results ("all results come from...")
methodological	A capability of the apparatus or analysis
synthesis	Integrates multiple results into one statement
interpretation	Reframes results through a theoretical lens
literature-context	A cited prior result the paper depends on

Getting the role right determines where the claim sits in the argument structure and what edges it can carry in the dependency graph.

How do we measure quality?

We ran the pipeline on **10 eLife papers** and compared the output against a **reference set** — an earlier iteration of the same extraction process, done semi-manually.

This is a **self-consistency metric**, not a ground truth comparison. It tells us: does the automated pipeline produce output that agrees with a prior analysis of the same papers?

Three metrics:

- **Recovery** — did the system find the same claims as the reference?
- **Role** — did it assign the same role (hypothesis, empirical, control, etc.)?
- **Panel** — did it assign the same figure panel?

Self-consistency results (10 papers)

METRIC	MEAN	TARGET	STATUS
Claim recovery	96.9%	80%	PASS
Role agreement	83.1%	75%	PASS
Panel agreement	60.9%	90%	FAIL

The system is **stable and reproducible** — it consistently finds the same claims and assigns similar roles. The external reviewer is what drives role agreement (without it: ~65%).

Panel agreement is a known limitation. Multi-panel claims collapse. Next iteration target.

Toward real ground truth

Self-consistency tells us the system is stable. These would tell us it's **correct**:

- **Human editorial benchmark** — editors or trained analysts create independent ground truth for a test set
- **Compare to peer reviewers** — digest review reports into claims, compare to extracted claims. Are they talking about the same things?
- **Compare to QED Science** — QED does claim-level structuring at journals. How do the two approaches compare?
- **Author validation** — show authors their paper's extracted claims. Do they agree?

Each gives a different angle on quality. These are natural next steps for a collaboration.

The system improves empirically

Every prompt change is tested against the 10-paper corpus before shipping.

Worked example: Kammer 2026 (a foveal feedback paper) scored 57% on role agreement — worst in the corpus. We looked at the 9 mismatches and found 3 patterns:

- The system labeled **control experiments** as plain `empirical` — it didn't recognize their eliminative function
- It called **hypotheses** `prediction` — shifting the argument layer down
- It missed **literature-context** claims — prior results the paper depends on but didn't produce

We added rules to the external reviewer prompt to catch each pattern, re-ran on that paper (cost: ~\$10), and role agreement jumped from 57% to 71%. Then we re-ran all 10 papers (~\$100): the corpus-wide mean rose from 78% to 83%, with only one small regression.

Prompt engineering as hypothesis testing, not guesswork.

Figure reproduction

Claims aren't just text — they link to **deposited code and data**. For 6 of the 10 papers, we ran the authors' analysis scripts and reproduced the figures.

- **68 of 245 claims** (28%) are backed by a reproduced figure
- **15 figures** reproduced across 6 papers — original and reproduction side by side on the website
- Reproductions run the authors' deposited code on their deposited data — no re-implementation

This is the verification layer: each claim links to its computation, and the computation has been independently executed. When a reproduced figure matches the original, the claim is **verified by code**.

The other 72% of claims are not yet verified — some papers lack deposited code, some analyses need environments we haven't set up. Expanding verification coverage is a natural next step.

Current status

What works:

- Full 8-step pipeline, automated end to end
- 10-paper validation with measured performance
- 32-page documentation site
- Version-controlled prompts, evaluation harness

What's still iterating:

- Panel assignment (60%, target 90%) — next prompt iteration target
- 2 of 10 papers below role threshold — methods papers need prompt variants
- Dependency mapping (Step 6) — currently analyst work
- Authoring tools — not started

This is a working draft, not a finished product. Iteration is part of the design.

Demo — the website

Let's look at a paper:

```
zmainen.github.io/elifе-claim-trees/
```

What you'll see for each paper:

- **Claim list** — all extracted claims with roles and panel anchors
- **Dependency graph** — interactive DAG showing how claims relate
- **Claim structure** — hierarchical view (hypothesis → predictions → tests)
- **Abstract alignment** — which claims map to which sentences in the abstract

Where could eLife plug in?

Recall the two paths from the introduction: **extract existing papers** and **build author tools**. The extraction pipeline is built. Here's where joint work would make the most difference:

- **Integration** — the pipeline currently outputs files. Wrapping it as a service for eLife's editorial system is straightforward engineering (~2 weeks for an HTTP API).
- **User experience** — we have a per-paper claim viewer. Cross-paper search, an editorial review surface, author-facing tools — these need design and product thinking that eLife's team is better positioned to do.
- **Improving the extraction** — the prompts are markdown files in version control. The evaluation harness measures every change. A data science team can iterate on prompt variants, test cheaper models, and expand the reference corpus.
- **Scaling** — at \$7/paper and 15 min, it works for hundreds of papers. For thousands, we'd need batch processing, concurrent execution, and domain-specific variants.

Economics at scale

SCALE	COST	TIME
1 paper	\$7	15 min
100 papers	\$700	~1 day (parallelized)
1,000 papers	\$7K	~3 days (parallelized)

Cost-reduction levers exist (batch APIs, cheaper models for some agents, concurrent execution) that could cut cost roughly in half — but they need engineering and quality measurement before deployment.

eLife scale (hundreds/year) is tractable today.

How to try it yourself

```
pip install elife-extract

# Extract claims from any eLife paper by DOI
elife-extract extract --doi 10.7554/eLife.95562 --output-dir ./out

# Review and write claim files
elife-extract write --draft ./out/draft-*.json \
  --corpus-dir ./claims --review-mode external

# See what you got
ls ./claims/headley-2024-*/
```

What you need: Python 3.10+, an Anthropic API key (or Vertex AI credentials).

Per paper: ~\$7, ~15 minutes. The 10 papers in the current corpus were selected by eLife — you could run more of your own the same way.

Full walkthrough: [docs/cli/first-paper/](#)

The bigger picture — claim-native publication

Three layers of work, each a step toward science where claims are the fundamental unit:

LAYER	WHAT	STATUS
1. Extraction	Convert existing papers to claim-level structure. The pipeline this talk is about.	Built and working
2. Authoring	Tools that help authors compose papers as claim-units from the start. Makes extraction less necessary over time.	Sketched; not built
3. Sandbox	A controlled environment where AI agents publish at claim-level granularity — an experiment in what scientific knowledge production looks like at the right unit.	Vision

The pipeline is layer 1 — it exists because the existing corpus needs handling. Layer 2 reduces that need. Layer 3 is what becomes possible when claim-native is the default.

Next steps

- **Try more papers** — pick papers from different domains and structures. The current 10 are neuroscience; how does it handle other fields?
- **Establish ground truth** — editorial benchmarks, reviewer comparison, author validation
- **Improve the extraction** — panel assignment, methods-paper variants, edge mapping
- **Build the experience** — editorial review surfaces, cross-paper browsing, author tools

Documentation

Docs site: `zmainen.github.io/elifе-claim-trees/docs/`

Covers: methodology, architecture, CLI usage, validation results, cost projections, API reference, contributor guide.

Source repo: Python CLI, version-controlled prompts, evaluation harness.

Contact: Zach Mainen, zmainen@neuro.fchampalimaud.org